

# Cellula Technologies

## – Task(5)

- Internship 1 – Group 5
- Malak Ahmed Saber
- Marawan Abbas Mohamed

# Data Overview

- 500000 row x 26 features about Drivers names, fare amount, weather condition, traffic condition and more
- 9 columns with 5 common null values, filled with their median
- 0 Duplicates
- Outliers :
  - Detected with boxplots
  - Removed with std dev method / used also Z\_score as another method

```
[120]: for col in df.select_dtypes(include='number').columns:  
        mean = df[col].mean()  
        stddev = df[col].std()  
        df = df[(df[col] >= mean - 3 * stddev) & (df[col] <= mean + 3 * stddev)]
```

# Feature selection

...	Feature	VIF
0	Weather_cloudy	inf
1	Weather_rainy	inf
2	Weather_stormy	inf
3	Weather_sunny	inf
4	Weather_windy	inf
5	Traffic Condition_Congested Traffic	inf
6	Traffic Condition_Dense Traffic	inf
7	Traffic Condition_Flow Traffic	inf
8	Car Condition	1.00
9	pickup_longitude	407.94
10	pickup_latitude	226.53
11	dropoff_longitude	472.51
12	dropoff_latitude	282.40
13	passenger_count	1.00
14	hour	1.02
15	day	1.00
16	month	1.01
17	weekday	1.02
18	year	1.01
19	jfk_dist	131.73
20	ewr_dist	1626.02
21	lga_dist	181.69
22	sol_dist	911.13
23	nyc_dist	203.74
24	distance	1.47
25	bearing	1.68

Making sure features aren't strongly related by Checking multicollinearity

	features	importance
16	distance	0.733279
0	pca0	0.051372
1	pca1	0.040228
15	year	0.039447
17	bearing	0.034248
2	pca2	0.023259
11	hour	0.021657
12	day	0.013510
13	month	0.012565
14	weekday	0.009888
9	Car Condition	0.004833
10	passenger_count	0.004675
5	Weather_stormy	0.001897
7	Traffic Condition_Congested Traffic	0.001887
8	Traffic Condition_Dense Traffic	0.001883
3	Weather_cloudy	0.001822
6	Weather_sunny	0.001785
4	Weather_rainy	0.001765

Used RandomForestRegressor to find the most important features.

distance	0.785844
pca2	0.479855
year	0.462017
pca1	0.090528
bearing	0.058874
pca0	0.054168
hour	0.014831
month	0.012855
passenger_count	0.008392
weekday	0.002283
day	0.001853

Checked mutual information as an extra measure for feature importance and take largest 10.

# Feature Engineering

## Categorical Columns:

- Dropped : 'User ID', 'User Name', 'key', 'pickup\_datetime'
- Label Encoding : 'Car Condition', 'Weather', 'Traffic Condition'
- Frequency Encoding : 'Driver Name' then dropped the original column

## Numerical Columns:

- PCA Applied on : 'pickup\_longitude', 'pickup\_latitude', 'dropoff\_longitude', 'dropoff\_latitude', 'jfk\_dist', 'ewr\_dist', 'lga\_dist', 'sol\_dist', 'nyc\_dist'
- PCA applied due to high multicollinearity

# Train Test Split

Target  
'fare amount'

Train : Test  
80% : 20%

Random State  
42

# Modelling

- Models Used :
  - Random Forest Regressor
  - Linear Regression
  - KNN
  - XGBOOST
  - Ridge Regression
  - Decision Tree

# Models Evaluation

## Random Forest

```
Test Set:
MAE: 1.364682667703858
RMSE: 2.2846530280952537
R²: 0.8219737503509096
```

```
Train Set:
MAE: 0.9201935920317754
RMSE: 1.5726088798368119
R²: 0.914680667937954
```

## Linear Regression

```
Test Set:
MAE: 1.8131996212865702
RMSE: 2.8079922582701884
R²: 0.7310724164261337
```

```
Train Set:
MAE: 1.7887366282621222
RMSE: 2.740245241342703
R²: 0.7409492439747407
```

## XGBOOST

```
Test Set:
MAE: 1.3183865309043195
RMSE: 2.241493292549924
R²: 0.8286364602178582
```

```
Train Set:
MAE: 1.2528237208896473
RMSE: 2.0503375468662712
R²: 0.8549703732666851
```

## KNN

```
Test Set:
MAE: 1.7528385575620018
RMSE: 2.6910778615036537
R²: 0.7530005055706375
```

```
Train Set:
MAE: 1.4170319572067602
RMSE: 2.135453495195475
R²: 0.842679167589044
```

## Ridge Regression

```
Test Set Evaluation:
MAE: 1.7922271830177852
RMSE: 2.7242017405769805
R²: 0.7402845311982003
```

```
Train Set Evaluation:
MAE: 1.8000608473367579
RMSE: 2.750899402526535
R²: 0.7392918562660663
```

## Decision Tree

```
Test Set Evaluation:
MAE: 1.7127912953065116
RMSE: 2.5923888872894243
R²: 0.7648096018698163
```

```
Train Set Evaluation:
MAE: 1.7119486811496598
RMSE: 2.610386922325979
R²: 0.7652449467477476
```

# Hyperparameter Tuning

- Used random grid for faster tuning
- After that, adjustments were made manually to ensure maximum efficiency.

```
from sklearn.model_selection import RandomizedSearchCV
xgb_model = XGBRegressor(n_estimators=300)

param_grid = {
    'learning_rate' : [0.01, .03, 0.1],
    'max_depth': [6, 10, 15],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'gamma': [0.1, 0.3, 1],
    'reg_alpha': [0, 0.1, 1.0],
    'reg_lambda': [1.0, 2.5, 5.0],
}

random_search = RandomizedSearchCV(
    estimator=xgb_model,
    param_distributions=param_grid,
    n_iter=100,
    scoring='neg_root_mean_squared_error',
    cv=3,
    verbose=3,
    n_jobs=-1
)

random_search.fit(x_train_scaled, y_train)
```





# Best & Evaluation

## XGBOOST

```
model2 = XGBRegressor( learning_rate= 0.1, n_estimators= 300, subsample= 1.0, reg_lambda= 5.0,  
                        reg_alpha= 1, max_depth= 10, gamma= 1, colsample_bytree= 0.8)
```

```
Test Set:  
MAE: 1.2735375724645712  
RMSE: 2.192128773412807  
R2: 0.8361012420282169
```

```
Train Set:  
MAE: 1.0451376053631476  
RMSE: 1.687677888240604  
R2: 0.9017381083957117
```

## Random Forest Grid Search cv2



```
Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}  
MAE: 1.8885532  
RMSE: 3.02  
R2 Score: 0.6572
```